

Microservice Architecture with ASP.NET Core

Agenda

- Introduction to .NET Core
- .NET 8.0 SDK
- Modern C# and What's New
- ASP.NET Core Application Architecture
- Models
- Controllers
- Request Routing
- Web APIs
- Data Validation
- Error Handling

Microservice Architecture with ASP.NET Core


Agenda


- Logging
- Microservice Architecture
- Designing Microservices
- Integration Patterns
- Scale and Resilience
- Containerization
- Security


Microservice Architecture with ASP.NET Core

Introduction

- Evolution of the .NET Platform
- .NET SDKs and Runtimes
- Visual Studio and Visual Studio Code



- 
- 2015 .NET 4.6, C# 6.0, Mono 4.0, Visual Studio 2015, Visual Studio Code 1.0
Expression Bodied Members, Null Propagator, String Interpolation
 - 2016 Xamarin Acquisition, .NET Core 1.0, .NET Standard 1.0
Entity Framework Core 1.0
 - 2017 .NET 4.7, .NET Core 2.0, C# 7.0, Visual Studio 2017
ASP.NET Razor Pages, Out Variables, Tuples, Ref Locals and Returns
 - 2018 GitHub Acquisition, .NET Standard 2.0
Blazor Server
 - 2019 .NET 4.8, .NET Core 3.0, C# 8.0, Visual Studio 2019
gRPC, Default Interface Methods, Using Declarations, Nullable Reference Types
 - 2020 .NET 5, C# 9.0
Blazor WebAssembly, Records, Init Only Setters, Top-Level Statements
 - 2021 .NET 6, C# 10.0, Visual Studio 2022
.NET MAUI
 - 2022 .NET 7, C# 11.0
Standard-Term Support (STS) release primarily focused on performance

- 
- 2023 .NET 8, C# 12
Collection Expressions, ref readonly Parameters, Blazor Render Modes

Introduction

Evolution of the .NET Platform

- In 2016, Microsoft introduced a new variant of .NET called .NET Core
- Many components were completely rewritten
- Fully supported on Windows, macOS, and Linux
- Included a subset of the functionality provided by .NET Framework
 - Focused on web-based workloads (web UIs and services)
- Merged MVC and Web API into the core framework

Introduction

Evolution of the .NET Platform

- The version of .NET Core after 3.1 became the "main line" for .NET and was labeled .NET 5.0
- The ASP.NET framework in .NET still includes the name "Core" to avoid confusion with previous versions of ASP.NET MVC

Introduction

Evolution of the .NET Platform

- The entire .NET platform is made available as open-source
- Community contributions are encouraged via pull requests
 - Thoroughly reviewed and tightly controlled by Microsoft

github.com/dotnet

Introduction

.NET SDKs and Runtimes

- .NET Runtime
 - Different version for each platform
 - Provides assembly loading, garbage collection, JIT compilation of IL code, and other runtime services
 - Includes the dotnet tool for launching applications
- ASP.NET Core Runtime
 - Includes additional packages for running ASP.NET Core applications
 - Reduces the number of packages that you need to deploy with your application

Introduction

.NET SDKs and Runtimes

- .NET SDK
 - Includes the .NET runtime for the platform
 - Additional command-line tools for compiling, testing, and publishing applications
 - Contains everything needed to develop .NET applications (with the help of a text editor)

Introduction

.NET SDKs and Runtimes

- Each version of .NET has a lifecycle status
 - Standard Term Support (STS) – Includes the latest features with a support period of 18 months
 - Long Term Support (LTS) – Has an extended support period of three years
 - Preview – Not supported for production use
 - Out of support – No longer supported

dotnet.microsoft.com/download

Introduction

Visual Studio and Visual Studio Code

- Visual Studio is available for Windows and macOS
 - Full-featured IDE
- Visual Studio Code is available for Windows, macOS, and Linux
 - Includes IntelliSense and debugging features
 - Thousands of extensions are available for additional functionality

visualstudio.microsoft.com

Introduction

Visual Studio and Visual Studio Code

- JetBrains also offers an IDE for .NET development called Rider
- Available for Windows, macOS, and Linux
- Includes advanced capabilities in the areas of refactoring, unit testing, and low-level debugging

www.jetbrains.com/rider

Microservice Architecture with ASP.NET Core

.NET SDK

- Installation
- Version Management
- Command-Line Interface (CLI)

.NET SDK

Installation

- The .NET SDK is distributed using each supported platform's native install mechanism
- Requires administrative privileges to install
- A list of installed SDK versions is available by using the .NET Command Line Interface (CLI)

```
dotnet --list-sdks
```

- A complete list of all installed runtimes and SDKs (as well as the default version) is also available

```
dotnet --info
```


.NET SDK

Version Management

- By default, CLI commands use the newest installed version of the SDK
 - This behavior can be overridden with a global.json file

```
{  
  "sdk": {  
    "version": "6.0.14"  
  }  
}
```

- Will be in effect for that directory and all sub-directories

.NET SDK

Version Management

- Use of global.json files can allow developers to experiment with newer versions of the SDK while ensuring consistency for specific projects
- Include a global.json file in a source control repository to ensure every member of the team is using the same version of the SDK
 - Will generate an error if the specified SDK version is not present on the system

.NET SDK

Version Management

- While the SDK version (tooling) is specified using a global.json file, the runtime version is specified within the project file

```
<PropertyGroup>  
  <TargetFramework>net8.0</TargetFramework>  
</PropertyGroup>
```

.NET SDK

Version Management

- When an application is launched, it will automatically use the newest available runtime with the same major and minor version number
 - For example, if version 8.0 is specified, the application will automatically use the 8.0.3 runtime but will not automatically use version 8.1 of the runtime
- Allows for system administrators to apply security patches and runtime bug fixes without the need to recompile and re-deploy the application
- Behavior can be overridden by specifying a RollForward policy value

.NET SDK

Version Management

- The target framework for a project can be an older version than the version of the SDK that you are using
 - For example, you can use version 8 of the SDK to build an application that targets the .NET 6 runtime

```
<PropertyGroup>  
  <TargetFramework>net6.0</TargetFramework>  
</PropertyGroup>
```

- Recommended approach – Use the newest version of the tools possible and choose a runtime target based on your deployment environment

.NET SDK

Command-Line Interface (CLI)

- Many higher-level tools and IDEs use the CLI "under-the-covers"
- CLI commands consist of the driver ("dotnet"), followed by a "verb" and then possibly some arguments and options

.NET SDK

Command-Line Interface (CLI)

- dotnet new
 - Create a new project from an available template
- dotnet restore
 - Restore the dependencies for a project (download missing NuGet packages)
- dotnet build
 - Build a project and all its dependencies
- dotnet run
 - Run an application from its source code (performs a build if necessary)

.NET SDK

Command-Line Interface (CLI)

- dotnet test
 - Execute unit tests for a project
- dotnet publish
 - Pack an application and its dependencies into a folder for deployment
- And many more...

Lab

.NET SDK

- Create and run a .NET 8 console application using the CLI
- Create and run an ASP.NET Core application using the CLI

Microservice Architecture with ASP.NET Core

Modern C# and What's New

- Introduction
- Global Using Directives
- File-Scoped Namespace Declarations
- Top-Level Statements
- Nullable Reference Types
- Init Only Setters
- Record Types
- Deferred Execution

Modern C# and What's New

Introduction

- C# 9 introduced with .NET 5
- C# 10 introduced with .NET 6
- C# 11 introduced with .NET 7
- C# 12 introduced with .NET 8

Modern C# and What's New

Global Using Directives

- C# 10 introduced global using directive support
- If the global keyword is present, the using directive will be in effect for every file in the project

```
global using EComm.Core;
```

- Can be in any file but a good practice is to have a separate cs file for all the project's global using directives

Modern C# and What's New

Global Using Directives

- Starting with .NET 6, global using directives for common system namespaces can be included implicitly via a project setting

```
<ImplicitUsings>enable</ImplicitUsings>
```

- This setting is included in new projects by default
- For an ASP.NET project, there are a total of 16 namespaces that will be implicitly referenced

Modern C# and What's New

File-Scoped Namespace Declarations

- Typically, code within a namespace is defined within curly braces

```
namespace Acme.Models  
{  
    ...  
}
```

- C# now allows for a namespace declaration to specify that all the code within a file belongs to a specific namespace

```
namespace Acme.Models;  
...
```

Modern C# and What's New

Top-Level Statements

- A .NET application requires an entry point function named Main defined within a static class

```
class Program {  
    static void Main(string[] args) {  
        Console.WriteLine("Hello, World!");  
    }  
}
```

- The C# compiler (10 and later) can recognize executable code that is outside of a class as the code for the entry point and generate the necessary function and static class for you

```
Console.WriteLine("Hello, World!");
```

Modern C# and What's New

Top-Level Statements

- ASP.NET project templates combine the implicit using feature with top-level statements to minimize the amount of code required in Program.cs

```
var builder = WebApplication.CreateBuilder(args);  
var app = builder.Build();  
  
app.MapGet("/", () => "Hello World!");  
  
app.Run();
```


Modern C# and What's New

Nullable Reference Types

- By default, value types in .NET cannot be set to null
 - A variable can be defined as a nullable value type so that it can store a null value

```
int? num = null;
```

- Reference types can store null and default to null if not provided with an initial value

```
Product p; // p is null
```

Modern C# and What's New

Nullable Reference Types

- The most common exception encountered during .NET development is the `NullReferenceException`
 - Occurs when attempting to access the member of an object that is null
- Safety can be significantly improved by using types that cannot be null unless explicitly identified to allow it

Modern C# and What's New

Nullable Reference Types

- C# 8 introduced the idea of nullable reference types
 - Like values types, reference types are not allowed to be null unless the variable is defined as nullable
- Because of the impact on existing code, this feature was not enabled by default
- Could be enabled via the Nullable annotation in the project file

```
<Nullable>enable</Nullable>
```

- Starting with .NET 6, nullable reference types are enabled by default

Modern C# and What's New

Nullable Reference Types

- If enabled, compiler warnings will be generated when...
 - Setting a non-nullable reference type to null
 - Defining a reference type that does not initialize all non-nullable reference type members as part of construction
 - Dereferencing a possible null reference without checking for null (or using the null-forgiving operator)

```
string fn = person!.FirstName;
```

Modern C# and What's New

Nullable Reference Types

- It is a good idea to enable nullable reference types for new projects
- Refactoring an existing application to use nullable reference types could require a significant amount of effort

Modern C# and What's New

Init Only Setters

- It is very convenient to initialize the properties of an object by using object initialization syntax

```
var product = new Product { Name = "Bread", Price = 2.50 }
```

- However, in the past, this was only possible by defining the properties as writable

Modern C# and What's New

Init Only Setters

- It is now possible to define properties with init only setters
- Properties can be set as part of object initialization but become read-only after that

```
public class Product
{
    public string Name { get; init; }
    public double Price { get; init; }
}
```

Modern C# and What's New

Record Types

- Every type in .NET is either a value type or a reference type
 - Struct is a value type
 - Class is a reference type
- Value types are recommended to be defined as immutable and are copied on assignment
 - Use value semantics for equality
 - Supports additional safety and optimizations especially for concurrent programming with shared data

Modern C# and What's New

Record Types

- The record type introduced in C# 9 allows you to easily define an immutable reference type that supports value semantics for equality

```
public record Person
{
    public string LastName { get; }
    public string FirstName { get; }

    public Person(string first, string last) =>
        (FirstName, LastName) = (first, last);
}
```

Modern C# and What's New

Record Types

- None of the properties of a record can be modified once it's created
- Records do support inheritance
- It is easy to create a new record from an existing one via the with keyword

```
var person = new Person("Joe", "Smith");
Person brother = person with { FirstName = "Bill" };
```

Modern C# and What's New

Record Types

- Record types can be a very good fit for things like ViewModels and Data Transfer Objects (DTOs)

Modern C# and What's New

Deferred Execution

- When using LINQ methods in C#, it is important to be aware of deferred execution

Microservice Architecture with ASP.NET Core

ASP.NET Core Application Architecture

- Introduction
- NuGet Packages
- Application Startup
- Hosting Environments
- Middleware and the Request Pipeline
- Services and Dependency Injection

ASP.NET Core Application Architecture

Introduction

- Single stack for Web UI and Web APIs
- Modular architecture distributed as NuGet packages
- Flexible, environment-based configuration
- Built-in dependency injection support
- Support for using an MVC-based architecture or a more page-focused architecture by using Razor Pages
- Blazor allows for the implementation of client-side functionality using .NET code

ASP.NET Core Application Architecture

NuGet Packages

- NuGet is a package manager for .NET
 - www.nuget.org
- All the libraries that make up .NET (and many 3rd-party libraries) are distributed as NuGet packages
- NuGet package dependencies are stored in the project file

```
<PackageReference Include="Microsoft.EntityFrameworkCore" Version="8.0.3" />
```

ASP.NET Core Application Architecture

NuGet Packages

- The dotnet restore command will fetch any referenced NuGet packages that are not available locally
- Uses nuget.org as the package source by default
- Additional or alternative package sources (remote or local) can be specified by using a nuget.config file

ASP.NET Core Application Architecture

NuGet Metapackages

- Metapackages are a NuGet convention for describing a set of packages that are meaningful together
- Every .NET Core project implicitly references the Microsoft.NETCore.App package
 - ASP.NET Core projects also reference the Microsoft.AspNetCore.App package
- These two metapackages are included as part of the runtime package store
 - Available anywhere the runtime is installed

ASP.NET Core Application Architecture

Application Startup

- When an ASP.NET Core application is launched, the first code executed is the application's Main method
 - Generated by the compiler if using top-level statements
- Code in the Main method is used to...
 - Create a WebApplication object
 - Configure application services
 - Configure the request processing pipeline
 - Run the application

ASP.NET Core Application Architecture

Application Startup

- WebApplication's CreateBuilder method is typically used to create the WebApplicationBuilder object
- When the WebApplicationBuilder is created, it loads configuration information from...
 - appsettings.json and appsettings.{Environment}.json
 - User secrets (when running in Development environment)
 - Environment variables and command-line arguments

ASP.NET Core Application Architecture

Application Startup

- After the WebApplicationBuilder has been initialized, application services can be added
- WebApplicationBuilder's Build method is used to construct the WebApplication object and initialize the dependency injection system
- The WebApplication object is used to configure the request processing pipeline

ASP.NET Core Application Architecture

Application Startup

- A collection of framework services are automatically registered with the dependency injection system
 - IHostApplicationLifetime
 - Used to handle post-startup and graceful shutdown tasks
 - IHostEnvironment / IWebHostEnvironment
 - Has many useful properties (ex. EnvironmentName)
 - ILoggerFactory
 - IServer
 - And many others...

ASP.NET Core Application Architecture

Application Startup

- The environment for local machine development can be set in the launchSettings.json file
 - Overrides values set in the system environment
 - Only used on the local development machine
 - Is not deployed
 - Can contain multiple profiles

ASP.NET Core Application Architecture

Application Startup

- By default, clients can use HTTP/2 when selected during the TLS handshake; otherwise, HTTP/1.1 is used
- Additional configuration options are described in the documentation

ASP.NET Core Application Architecture

Hosting Environments

- EnvironmentName property can be set to any value
- Framework-defined values include:
 - Development
 - Staging
 - Production (default if none specified)
- Typically set using the ASPNETCORE_ENVIRONMENT environment variable
- Can also be configured via launchSettings.json

ASP.NET Core Application Architecture

Middleware

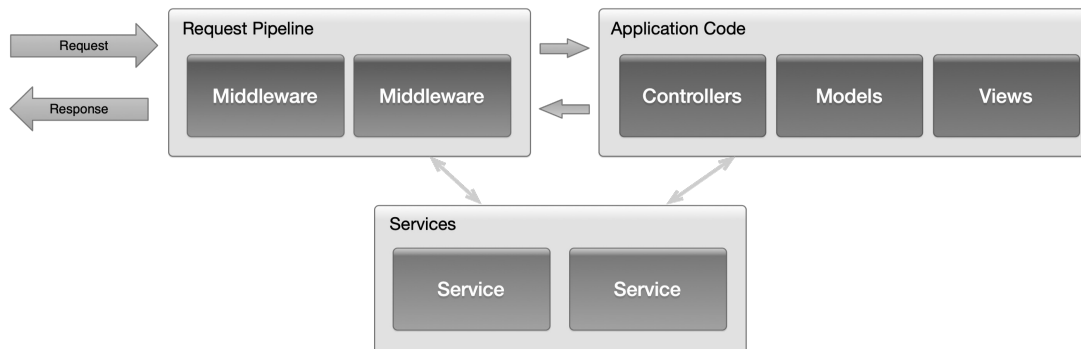
- ASP.NET uses a modular request processing pipeline
- The pipeline is composed of middleware components
- Each middleware component is responsible for invoking the next component in the pipeline or short-circuiting the chain
- Examples of middleware include...
 - Request routing
 - Handling of static files
 - Authentication
 - Response caching
 - Error handling

ASP.NET Core Application Architecture

Services

- ASP.NET Core also includes the concept of services
- Services are components that are available throughout an application via dependency injection
- An example of a service would be a component that accesses a database or sends an email message

ASP.NET Core Application Architecture



ASP.NET Core Application Architecture

Pipeline

- The last piece of middleware in the pipeline is typically the routing middleware
- Routes the incoming request to a controller
- Instead of using controllers, the new minimal API framework in .NET can be used

Microservice Architecture with ASP.NET Core

Models

- Introduction
- Persistence Ignorance
- Object-Relational Mapping
- Asynchronous Data Access

Models

Introduction

- Models represent "real world" objects the user is interacting with
- Entities are the objects used during Object-Relational Mapping and provide a way to obtain and persist model data
- The term Data Transfer Object (DTO) is often used to describe an object that carries data between different processes or subsystems
 - A single DTO may contain multiple different entities, exclude some entity properties, or use different property names
 - In a Web API application, the object that get serialized into JSON is often a DTO

Models

Persistence Ignorance

- The model data typically comes from an external source (database, web service, file, etc.)
- For better maintainability and testability, it is a best practice to use a data access component to encapsulate the details about where the model data comes from
- In ASP.NET, data access should be performed by a service made available via dependency injection
 - Makes it easy to test components independently with hard-coded data (no database)

Models

Object-Relational Mapping

- If a data access component communicates with a relational database, a necessary task will be to convert between relational data and C# objects
- This can be done manually by with ADO.NET, or several frameworks exist that can help with this task
 - Entity Framework Core
 - Dapper (3rd-party micro-ORM)
 - AutoMapper (mapping one object to another)

Models

Asynchronous Data Access

- When performing IO-bound operations (database access, web service calls, etc.), it is a best practice to perform that work asynchronously
- Allows for the efficient use of thread resources
 - Thread pool threads can be used to handle other incoming requests while the IO-bound operation is in progress
 - Improves the scalability of a web application

```
public async Task<IEnumerable<Product>> GetAllProducts()  
{  
    return await _repository.GetProductsAsync();  
}
```

Microservice Architecture with ASP.NET Core

Application Configuration

- Middleware
- Services
- Configuration Providers and Sources
- Configuration API
- Options Pattern

Application Configuration

Middleware

- A middleware component typically adds an extension method to `IApplicationBuilder` for adding it to the pipeline
 - By convention, these methods start with the prefix "Use"

```
app.UseHttpsRedirection();  
app.UseAuthorization();
```

- The order in which middleware is added to the pipeline can be important
 - Determines the order of execution
 - As an example, it would be very important for authentication middleware to execute before some caching middleware that could return a cached response

Application Configuration

Services

- Services are components that are available throughout an application via dependency injection
- The lifetime of a service can be...
 - Singleton (one instance per application)
 - Scoped (one instance per web request)
 - Transient (new instance each time component requested)
- An example of a service would be a component that accesses a database or sends an email message

Application Configuration

Services

- Services are typically added via extension methods available on IServiceCollection

```
builder.Services.AddDbContext<ApplicationDbContext>(...);  
builder.Services.AddScoped<IEmailSender, MyEmailSender>();  
builder.Services.AddScoped<ISmsSender, MySmsSender>();
```

- Most methods include the service lifetime as part of the method name (e.g., AddScoped)
- The AddDbContext method is a custom method specifically for adding an Entity Framework DbContext type as a service

Application Configuration

Services

- Services are available throughout the application via dependency injection
- A common practice is to follow the Explicit Dependencies Principle
 - Controllers include all required services as constructor parameters
 - System will provide an instance or throw an exception if the type cannot be resolved via the DI system

```
public class ProductController : ControllerBase  
{  
    public ProductController(IEmailSender emailSender) {  
        ...  
    }  
}
```

Application Configuration

Configuration Providers and Sources

- Before ASP.NET Core, application settings were typically stored in an application's web.config file
- ASP.NET Core introduced a completely new configuration infrastructure
 - Based on key-value pairs gathered by a collection of configuration providers that read from a variety of different configuration sources

Application Configuration

Configuration Providers and Sources

- Available configuration sources include:
 - Files (INI, JSON, and XML)
 - System environment variables
 - Command-line arguments
 - In-memory .NET objects
 - Azure Key Vault
 - Custom sources

Application Configuration

Configuration Providers and Sources

- The default WebApplicationBuilder adds providers to read settings (in the order shown) from:
 - appsettings.json
 - appsettings.{Environment}.json
 - User secrets
 - System environment variables
 - Command-line arguments
- Values read later override ones read earlier

Application Configuration

Configuration API

- The configuration API provides the ability to read from the constructed collection of name-value pairs
- An object of type IConfiguration is available to be used via dependency injection

```
public class HomeController : ControllerBase
{
    public HomeController(IConfiguration configuration)
    {
        _emailServer = configuration["EmailServer"];
    }
}
```

Application Configuration

Configuration API

- Hierarchical data is read as a single key with components separated by a colon

```
{
  "Email": {
    "Server": "gmail.com",
    "Username": "admin"
  }
}
```

```
public class HomeController
{
  public HomeController(IConfiguration configuration)
  {
    _emailServer = configuration["Email:Server"];
  }
}
```

Application Configuration

Options Pattern

- The options pattern can be used to provide configuration information to other components within your application as strongly-typed objects via dependency injection

```
public class EmailOptions
{
  public string Server { get; set; }
  public string Username { get; set; }
}
```

```
builder.Services.Configure<EmailOptions>(Configuration.GetSection("Email"));
```

```
public HomeController(IOptions<EmailOptions> emailOptions)
{
  _emailOptions = emailOptions;
}
```

Microservice Architecture with ASP.NET Core

Controllers

- Responsibilities
- Requirements and Conventions
- Dependencies
- Action Results
- Model Binding

Controllers

Responsibilities

- The action executed for a particular endpoint is typically a method of a controller
- A controller may need to retrieve or make modifications to model data
- The controller also determines the appropriate type of response to return
 - HTML, JSON, XML, redirection, error, etc.

Controllers

Responsibilities

- Controller methods that are reachable via the routing system are referred to as controller actions
- Any public method of a controller can be an action if a valid route to that action exists

Controllers

Requirements and Conventions

- For a class to act as a controller, it must...
 - Be defined as public
 - Have a name that ends with Controller or inherit from a class with a name that ends with Controller
- Common conventions (not requirements) are...
 - Place all controllers in a root-level folder named Controllers
 - Inherit from a system class called Controller (or its subclass ControllerBase for an API)
 - Provides many helpful properties and methods

Controllers

Dependencies

- It is a recommended best practice for controllers to follow the Explicit Dependencies Principle
- Specify required dependencies via constructor parameters that can be supplied via dependency injection

```
public class HomeController : Controller
{
    private IEmailSender _emailSender;

    public HomeController(IEmailSender es) {
        _emailSender = es;
    }
}
```

Controllers

Action Results

- ActionResult<T> enables returning a type derived from ActionResult or a specific type
- Framework uses the ExecuteResultAsync when creating the HTTP response

```
public ActionResult<Product> GetProduct(int id)
{
    var product = _repository.GetProduct(id);
    if (product == null) return NotFound();
    return product;
}
```

- Writing directly to the response should be avoided
 - Adds a dependency to the HTTP context
 - Make things like unit testing more difficult

Controllers

Action Results

- The base class Controller provides helper methods to generate various types of results
 - View `return View(customer);`
 - Serialized object `return Json(customer);`
 - HTTP status code `return NotFound();`
 - Raw content `return Content("Hello");`
 - Contents of a file `return File(bytes);`
 - Several forms of redirection
 - Redirect, RedirectToRoute, RedirectToAction, ...
 - And more...

Controllers

Asynchronous Controller Actions

- It is common for a controller action to invoke an asynchronous method to perform an IO-bound operation
 - Database access, web service call, etc.
- The action should be marked as async with a return type of Task<T> and await used with the asynchronous method

```
public async Task<IEnumerable<Product>> Index()
{
    var products = await _repository.GetAllProducts();
    return products;
}
```

Controllers

Asynchronous Controller Actions

- Making an action asynchronous does not change the experience for the client
 - No response is sent until the entire action is complete
- Can improve application scalability by allowing the thread pool thread to handle other incoming requests while waiting for the IO-bound operation to complete
- It is also possible to accept a CancellationToken that can be used to handle the cancellation of a long-running request

```
public async Task<IEnumerable<Product>> Index(CancellationToken token)
{
    var products = await _repository.GetAllProducts(token);
    return products;
}
```

Controllers

Model Binding

- When an action is invoked, the model binding system attempts to populate the parameters of the action with values in the request
 - Body values
 - Route value
 - Query strings
- Items above are listed in priority order (i.e., body values will take precedence over other values)

Controllers

Model Binding

- If an action accepts an object parameter, the model binding system will create an instance of that type and attempt to populate its public properties with values from the request
- If validation errors occur during the model binding process, the IsValid property of the ModelState property will return false

```
public ActionResult Edit(ProductViewModel vm)
{
    if (ModelState.IsValid) { ... }
    ...
}
```

Controllers

Model Binding

- It is important to ensure the model binding system does not alter values that you do not intend to be modified
 - Can lead to a security vulnerability known as over-posting
- Attributes can be used to define properties that should not participate in model binding

```
[BindNever]
public int EmployeeId { get; set; }
```

- Alternatively, use a DTO that only includes properties that are intended to participate in model binding

Microservice Architecture with ASP.NET Core

Request Routing

- RESTful Services
- Endpoint Routing
- Route Attributes
- Route Templates
- Route Constraints
- Route Template Precedence

Request Routing

RESTful Services

- When configuring request routing, you should try to maintain a RESTful API
- Clean, extension-less URLs that identify resources
- Use of the correct HTTP verbs within an API
- Avoid query string parameters except for ancillary data that is related to the presentation of the information
 - Sorting key, current page number, etc.

Request Routing

Endpoint Routing

- Routing is responsible for mapping request URIs to endpoints and dispatching incoming requests to those endpoints
- Routing can also be used to generate URLs that map to endpoints
 - Eliminates hardcoded URLs that would need to be updated when the routing configuration changes

Request Routing

Endpoint Routing

- A collection of extension methods on `ApplicationBuilder` are available for adding different types of endpoints
 - All start with the word "Map"
- Verb-based methods make it easy to configure simple endpoints
 - `MapGet`, `MapPost`, `MapPut`, `MapDelete`, etc.
- `MapControllers` will configure and add an endpoint for each controller action defined in the application

Request Routing

Route Attributes

- If using controllers, attributes can be used to define the routing information used to construct the endpoints
- The Route attribute will create an endpoint for all HTTP verbs

```
[Route("products")]  
public IActionResult AllProducts() { ... }
```

- Verb-specific attributes should be used to define an endpoint for a specific HTTP verb

```
[HttpGet("products")]  
public IActionResult AllProducts() { ... }
```

Request Routing

Route Attributes

- A controller-level attribute can be used to specify a prefix for all the actions of the controller

```
[Route("[controller]")]  
public class ProductController : Controller  
{  
    [HttpGet("{id}")]  
    public IActionResult GetProduct(int id) { ... }  
}
```

- In the example above, a request for a product would use a URL of <https://example.com/product/6>

Request Routing

Route Templates

- Tokens within curly braces define route value parameters which will be bound if the route is matched
 - You can define more than one route value parameter in a route segment, but they must be separated by a literal value

`site/{name}/{id}`

`{language}-{region}/library/{topic}`



`{language}{region}/{topic}`

Request Routing

Route Templates

- Route value parameters can have default values
 - The default value is used if no value is present in the URL for the parameter

`products/{sort=price}`

- Route value parameters may also be marked as optional
 - When bound to an action parameter, the value will be null (reference type) or zero (value type)

`product/{id?}`

Request Routing

Route Templates

- The catch-all parameter (identified using an asterisk) allows for a route to match a URL with an arbitrary number of parameters

```
query/{category}/{*path}
```

```
http://localhost/query/people/hr/managers
```

```
public IActionResult Query(string category, string path)
{
    // category = "people"
    // path = "hr/managers"
}
```

Request Routing

Route Constraints

- A route value parameter can include an inline constraint
- URLs that do not match the constraint are not considered a match
- Multiple constraints can be specified for one parameter

```
products/{id:int}
```

```
products/{id:range(100, 999)}
```

```
employees/{ssn:regex(d{3}-d{2}-d{4})}
```

```
products/{id:int:range(100, 999)}
```

Request Routing

Route Constraints (Partial List)

Constraint	Example Route	Example Match
int	{id:int}	123
bool	{active:bool}	true
datetime	{dob:datetime}	2016-01-01
guid	{id:guid}	7342570B-44E7-471C-A267-947DD2A35BF9
minlength(value)	{username:minlength(5)}	steve
length(min, max)	{filename:length(4, 16)}	Somefile.txt
min(value)	{age:min(18)}	19
max(value)	{age:max(120)}	91
range(min, max)	{age:range(18, 120)}	91
alpha	{name:alpha}	Steve
regex(expression)	{ssn:regex(d{3}-d{2}-d{4})}	123-45-6789

Request Routing

Route Constraints

- Route constraints should be used to help determine the route that should be used but should not be used for the validation of input values
- If a matching route is not found, the response from the server will be a 404 (resource not found)
- Invalid input should typically result in a different response (e.g., 400 with an appropriate error message)

Request Routing

Route Template Precedence

- Each route template is assigned a value by the system based on how specific it is
- Literal segments are considered more specific than parameter segments

`/hello/customer`

`/hello/{name}`

- A parameter segment with a constraint is considered more specific than one without a constraint
- The Order property of an endpoint can be used to override the default precedence behavior
- If a request matches multiple endpoints with the same precedence, an `AmbiguousMatchException` is thrown at runtime

Microservice Architecture with ASP.NET Core

Web APIs

- API Controllers
- OpenAPI / Swagger
- Testing APIs
- Retrieval Operations
- Model Binding
- Update, Create, and Delete Operations
- Cross-Origin Request Sharing (CORS)

Web APIs

API Controllers

- ASP.NET Core includes a class named ControllerBase
 - Includes many properties and methods for handling HTTP requests
- The Controller class inherits from ControllerBase and adds support for views
- If creating a controller that does not have any views, you should inherit directly from ControllerBase

Web APIs

API Controllers

- An API controller should be decorated with the ApiController attribute

```
[ApiController]  
public class ProductApiController : ControllerBase
```

- Automatic HTTP 400 responses for validation failures
- Problem details for error status codes

Web APIs

OpenAPI / Swagger

- OpenAPI is a specification for describing REST APIs
- Swagger is a collection of tools that work with OpenAPI
 - SwaggerDocument objects expose data about the API in JSON format (openapi.json)
 - Swagger UI is a dynamically generated web-based UI that can be used to view and test API methods

Web APIs

OpenAPI / Swagger

- By default, the API project templates include a reference to Swashbuckle.AspNetCore
- SwaggerDocument generation is handled by a service

```
services.AddSwaggerGen();
```

- Document availability and Swagger UI is configured via middleware components

```
app.UseSwagger();  
app.UseSwaggerUI();
```

Web APIs

OpenAPI / Swagger

- The ProducesResponseType attribute should be used when defining Web API actions

```
[HttpPost]  
[ProducesResponseType(StatusCodes.Status201Created)]  
[ProducesResponseType(StatusCodes.Status400BadRequest)]  
public ActionResult<Product> Create(Product product)
```

- Used by tools like Swagger to generate more descriptive documentation

Web APIs

OpenAPI / Swagger

- Actions (or entire controllers) can be omitted from the Swagger document generation process by using the ApiExplorerSettings attribute

```
[ApiExplorerSettings(IgnoreApi = true)]  
public class ErrorController : Controller
```

Web APIs

Testing APIs

- API endpoints that are exposed via GET are easy to test using a web browser
- For other verbs, it can be helpful to have a tool that can be used to craft custom HTTP requests
 - Postman application is very popular (getpostman.com)
 - Many other options are available

Web APIs

Testing APIs

- Microsoft introduced a new tool called the HTTP REPL

```
dotnet tool install -g Microsoft.dotnet-httprepl
```

- Command-line tool for making HTTP requests
- Supports most of the HTTP verbs
- Can use Swagger documents to discover the endpoints

```
> https://localhost:5001/~ ls
```

```
Products [get|post]  
Customers [get|post]
```

[docs.microsoft.com/en-us/aspnet/core/
web-api/http-repl](https://docs.microsoft.com/en-us/aspnet/core/web-api/http-repl)

Web APIs

Retrieval Operations

- In a Web API, retrieval operations are performed with an HTTP GET request
- If successful, the response should use an HTTP 200 status code

Web APIs

Retrieval Operations

- There are several options available for altering the format of the JSON returned

- Attributes
- Custom formatter
- Data projection

```
[JsonPropertyName("price")]  
public double UnitPrice { get; set; }
```

```
public async Task<ActionResult> GetProduct(int id)  
{  
    var product = await _repository.GetProduct(id, true);  
    if (product == null) return NotFound();  
    var retVal = new {  
        Id = product.Id, Name = product.ProductName,  
        Price = product.UnitPrice,  
        Supplier = product.Supplier.CompanyName  
    };  
    return Ok(retVal);  
}
```


Web APIs

Create Operations

- In a Web API, create operations are performed with an HTTP POST request
- If successful, the response should use an HTTP 201 (created) status code with a Location header set to the URI of the newly created resource
- The CreatedAtAction and CreatedAtRoute methods can be used to generate a correctly formatted response

```
return CreatedAtAction("GetProduct", new { id = product.Id }, product);
```

Web APIs

Delete Operations

- In a Web API, delete operations are performed with an HTTP DELETE request
- If successful, the response should use an HTTP 204 (no content) status code

```
return new NoContentResult();
```

Web APIs

Update Operations

- In a Web API, update operations are performed with...
 - HTTP PUT – Replaces an existing resource
 - HTTP PATCH – Modifies part of an existing resource
- If successful, the response should be HTTP 204 (no content)

```
public async Task<ActionResult> PutProduct(int id, Product product)
{
    if (id != product.Id) return BadRequest();

    var existingProduct = await _repository.GetProduct(id);
    if (existingProduct == null) return NotFound();

    await _repository.SaveProduct(product);
    return NoContent();
}
```

Web APIs

Cross-Origin Resource Sharing (CORS)

- Browser security prevents a web page from making Ajax requests to another domain
- CORS is a W3C standard that allows a server to relax this policy
- A server can explicitly allow some cross-origin requests
- CORS is configured in ASP.NET Core via a service and middleware

```
builder.Services.AddCors(options => {
    options.AddPolicy("MyCorsPolicy",
        policy => {
            policy.WithOrigins("https://myapp.com");
        });
});
```

```
app.UseCors("MyCorsPolicy");
```

Microservice Architecture with ASP.NET Core

Data Validation

- Introduction
- Data Annotations
- Model Binding
- IValidatableObject

Data Validation

Introduction

- Whenever any data from the client is being used to perform an action, it is important to have data validation in place
 - Don't skip validation for hidden form fields, HTTP header values, cookies, etc. (all are easy to modify)
- Client-side validation provides a good user experience and improved application scalability (less trips to the server)
- Server-side validation must also be provided
 - Client-side validation is easy to circumvent or may not be supported on the client

Data Validation

Data Annotations

- A variety of data annotations can be added to the model (or view model) that is sent to a view
- Data annotations are used during model binding to perform server-side validation

```
public class ProductEditViewModel
{
    [Required]
    public string ProductName { get; set; }
```

Data Validation

Data Annotations

Attribute	Purpose
[Required]	Property value is required (cannot allow nulls)
[StringLength]	Specifies a maximum length for a string property
[Range]	Property value must fall within the given range
[RegularExpression]	Property value must match the specified expression
[Compare]	Property value must match the value of another property
[EmailAddress]	Property value must match the format of an email address
[Phone]	Property value must match the format of a phone number
[Url]	Property value must match the format of a URL
[CreditCard]	Property value must match the format of a credit card number

Data Validation

Model Binding

- If a value is considered to be invalid, an error is added to ModelState and ModelState.IsValid will return false
- ModelState is also used by the helpers in the view when returning a view after a server-side validation error

Data Validation

IValidatableObject

- For custom server-side validation, you can implement the IValidatableObject interface for the type being populated by the model binder
- Any errors returned are automatically added to ModelState by the model binder

```
public IEnumerable<ValidationResult> Validate(ValidationContext
                                             validationContext)
{
    var retVal = new List<ValidationResult>();
    if (BirthDate > HireDate) {
        retVal.Add(new ValidationResult("Employee cannot be
                                       hired before they were born"));
    }
    return retVal;
}
```

Microservice Architecture with ASP.NET Core

Error Handling

- Best Practices
- HTTP Error Status Codes

Error Handling

Best Practices

- Handle errors as best you can when they occur
- Record the error information and/or send a notification
- Provide the user with an appropriate response
 - Do not reveal information that a malicious user could potentially use against you (e.g., database schema information)
 - Give the user some options (e.g., link to visit the home page in the case of a 404)
 - Use static content whenever possible to avoid an error page that itself produces an error

Error Handling

HTTP Error Status Codes

- The HTTP protocol defines a range of status codes that signify an error
 - 4xx = client error (not found, bad request)
 - 5xx = server error
- It is a best practice to define an appropriate customized response that will be returned to the client in these cases

Microservice Architecture with ASP.NET Core

Logging

- Introduction
- Configuration
- ILogger
- Serilog and Seq

Logging

Introduction

- Just as important as error handling is the ability to record information about events that occur
- Logging of error information is essential for tracking down an issue that occurs in production
- It is sometimes helpful to record information about events that are not errors
 - Performance metrics
 - Authentication audit logs

Logging

Introduction

- ASP.NET Core has a logging API that works with a variety of logging providers
- Built-in providers allow you to log to the console and the Visual Studio Debug window
- Other 3rd-party logging frameworks can be used to provide many other logging options
 - Serilog
 - NLog
 - Log4Net
 - Loggr
 - elmah.io

Logging

ILogger

- Any component that wants to use logging can request an `ILogger<T>` as a dependency

```
public class ProductController : Controller
{
    public ProductController(IRepository repository,
        ILogger<ProductController> logger) { }
}
```

Logging

ILogger

- `ILogger` defines a set of extension methods for different verbosity levels
 - Trace (most detailed)
 - Debug
 - Information
 - Warning
 - Error
 - Critical

```
_logger.LogInformation("About to save department {0}", id);
```

Logging

ILogger

- The highest verbosity level written to the log is typically set in appsettings

```
"Logging": {  
  "LogLevel": {  
    "Default": "Debug",  
    "System": "Information",  
    "Microsoft": "Information"  
  }  
}
```

Logging

Serilog

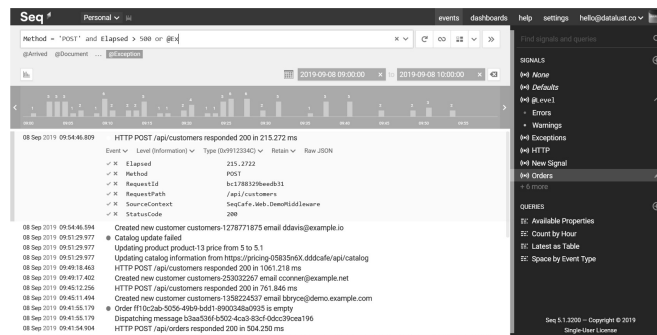
- Serilog has become a popular choice for ASP.NET Core
 - Wide variety of destinations and formats
 - Can record structured event data

github.com/serilog/serilog-aspnetcore

Logging

Seq

- In many ASP.NET Core applications, the log data needs to be off-host and centralized (e.g., load-balanced environment)
- Seq is an open-source server that can accept logs via HTTP
 - Integrates with .NET Core, Java, Node.js, Python, Ruby, Go, Docker, and more



© Treeloo, Inc. - All rights reserved (24-075)

133

Logging

Maintainability

- One important note is that the logging framework(s) you choose should not change how you write to the log (ILogger)
 - The only code that changes is in Program.cs

© Treeloo, Inc. - All rights reserved (24-075)

134

Microservice Architecture with ASP.NET Core

Microservice Architecture

- Introduction
- The Monolith
- Service Oriented Architecture
- Microservice Defined
- Service Boundaries
- Key Principals

Microservice Architecture

Introduction

- Building distributed systems is difficult
- Microservices can optimize for one kind of problem while making other parts of your system more complex
- A monolithic application on a single server will be simpler in many ways

Microservice Architecture

The Monolith

- Single process
- Single code base
- Shared database

Microservice Architecture

The Monolith

- A modular monolith can help with maintainability
- Well defined interfaces
- Unit tests

Microservice Architecture

The Monolith

- Where a modular monolith struggles is with scalability
- When the load on an application increases, the load tends to not be distributed evenly
- Scaling a monolith to handle more load can be difficult and wasteful
- Deployment coupling makes the system fragile
- Monoliths also make it difficult to scale a team of developers beyond a certain point

Microservice Architecture

Service Oriented Architecture

- Introduced as an architecture style in 1998
- First step toward breaking a monolith down into smaller pieces

Microservice Architecture

Service Oriented Architecture

- One service tends to be aligned with an entire business domain
 - Payment service, order service, inventory service, etc.
- All data exposed via service interfaces
- All communication via service calls over the network
- Presence of a network introduces additional complexity and failure points

Microservice Architecture

Service Oriented Architecture

- Presence of a network introduces additional complexity, latency, and failure points

Microservice Architecture

Microservices Defined

- An independently deployable unit that performs one job extremely well
- Interacts with other services via a clearly defined interface
- Done correctly, microservices can provide flexibility, scalability, and resilience
- Ideally, microservices should be able to evolve independently
 - Avoid breaking backward compatibility of interfaces

Microservice Architecture

Microservice Boundaries

- Defining microservice boundaries is a critical step in the design process
- Two tools are primarily used to help find boundaries
 - Domain Driven Design
 - Event Storming

Microservice Architecture

Event Storming

- Should involve technical and non-technical stakeholders
- Start with a list of all events that could occur
- For each...
 - Who triggered this event?
 - What command is given?
 - What action was taken?
 - What system is this part of?

Microservice Architecture

Key Principals

- Service Independence
 - Independently deployable and scalable
 - Can be developed and maintained by separate teams
- Decentralized Data
 - Each microservice should maintain and be in complete control of its own data store
 - Ensures loose coupling and autonomy

Microservice Architecture

Key Principals

- API Based Communication
 - Involves committing to a contract
 - Commonly REST APIs via HTTP
 - Can also use other mechanisms like gRPC or a message queue
- Resilience and Fault Tolerance
 - Microservices should be designed to handle failures gracefully

Microservice Architecture

Key Principals

- Scalability
 - Additional resources can be added to a single microservice
- Autonomous Teams
 - Enables faster development and innovation
 - Reduces blocking

Microservice Architecture with ASP.NET Core

Designing Microservices

- Coupling and Cohesion
- Sharing Data
- Synchronous Communication
- Asynchronous Communication
- Backward Compatibility
- Cross-Cutting Concerns
- Service Discovery

Designing Microservices

Coupling and Cohesion

- Coupling is the degree of interdependence between modules in your system
- Cohesion is the degree in which elements inside a module belong together
 - “Things that change together, live together”

Designing Microservices

Coupling and Cohesion

- A service boundary determines...
 - How your services interact
 - What they share
 - What they don't share
- If a change to one service requires changes to other services, you may have a distributed monolith

Designing Microservices

Sharing Data

- Each service owns its own data
- No services have direct access to the data of another
- Therefore, accessing shared data is typically the biggest area where coupling comes from

Designing Microservices

Sharing Data

- Information hiding is the principle of hiding as much data as possible behind a microservices boundary
- Data Transfer Objects (DTOs) are useful for providing an abstraction layer over the underlying data store
 - Allows the data schema to vary independently from the API

Designing Microservices

Synchronous Communication

- The most common form of communication between microservices is synchronous communication
 - For example, an HTTP request / response
- Immediate feedback for operation completion / failure
- Introduces runtime coupling
 - Service needs to know how to communicate with the external service
- Latency of external service adds to latency of the calling service
- Increases the chance of a cascading failure

Designing Microservices

Asynchronous Communication

- Allows a service to make a non-blocking call to another service
- Reducing runtime coupling
 - In a message-based architecture, the schema of a message is a dependency and an alternative form of runtime coupling
- Can add significant complexity
 - Operation status, completion, and timeouts
 - Eventual consistency

Designing Microservices

Eventual Consistency

- With asynchronous communication...
 - You cannot guarantee the state of another microservice
 - You need to handle failures differently
- Most microservice architectures end up including a mix of synchronous and asynchronous communication
- As a general best practice, start with asynchronous options and use synchronous communication when necessary

Designing Microservices

Backward Compatibility

- It is not possible to eliminate all forms of runtime coupling
- When changing an interface or message schema, always strive for maintaining backward compatibility
- You can offer alternate versions of a service

```
/api/v1/orders  
/api/v2/orders
```

```
{  
  "messageVersion": "v2",  
  "messageType": "orderCreated"  
  ...  
}
```

- Should be used with care and with a process to eventually remove older versions

Designing Microservices

Cross-Cutting Concerns

- Code sharing among microservices should be minimized
- Resist the urge to apply the DRY principle to microservices for domain logic
- Common libraries should only be used for generic cross-cutting concerns (e.g., logging)
 - Multiple microservices running different versions of the shared library should not present an issue

Designing Microservices

Service Discovery

- Less of a problem for asynchronous communication
 - One central message bus or queue
- For synchronous communication, service A needs a way to know how to communicate with service B
- DNS provides a simple form of service discovery
- An API Gateway maps external requests onto internal services
 - Can terminate things like HTTPS and authentication

Microservice Architecture with ASP.NET Core

Integration Patterns

- Introduction
- Shared Data
- Request / Response
- Publish / Subscribe
- Point to Point
- Distributed Transactions
- API First Design

Integration Patterns

Introduction

- "It's not the boxes you need to worry about in microservices, it's the lines"
- You can facilitate the integration of microservices by using enterprise integration patterns

Integration Patterns

Shared Data

- A shared data integration involves an agreed upon data transfer format and a shared storage location
- Form of asynchronous integration
- Useful for integration between disparate legacy systems but other options should be considered first

Integration Patterns

Request / Response

- Most common integration pattern in distributed systems
- Includes RPC, REST, and GraphQL
- RPC based services provide better performance at the cost of interoperability and tighter coupling
- GraphQL is a query language and an API runtime that allows clients to request specific data from servers
 - Allows clients to specify the structure of the data returned

Integration Patterns

Publish / Subscribe

- Messaging pattern that decouples publishers and subscribers
 - Some coupling does still exist at the message schema level
- A single message type could have multiple publishers and multiple subscribers
- AWS SNS is an example
- Observability, sequencing, and durability can present challenges

Integration Patterns

Point to Point

- Asynchronous form of communication but is one-to-one
- Can be durable and enforce sequencing
- AWS SQS is an example

Integration Patterns

Distributed Transactions

- Most developers are familiar with performing multiple database operations in the context of a single transaction
- With multiple microservice in control of their own data, this practice becomes more challenging
- One solution is to use a two-phase commit
 - Can introduce blocking and limit scalability

Integration Patterns

Distributed Transactions

- Another approach is the Saga pattern
- Saga is a collection of individual tasks
- State of the Saga often managed by a separate microservice
- Compensation (rollback) logic handled by the microservices responsible for each individual task

Integration Patterns

API First Design

- When starting a new integration project, agree on the API first
 - Could include message schemas
- Different teams can develop based on the agreed upon contract
- Start with a short period of synchronous collaboration followed by independent development processes

Microservice Architecture with ASP.NET Core

Scale and Resilience

- Introduction
- Timeouts and Retries
- Circuit Breakers
- Bulkheads
- Idempotency
- Vertical vs. Horizontal Scaling
- Caching
- Chaos Engineering

Scale and Resilience

Introduction

- Resilience
 - Robustness – Ability to accommodate expected failures
 - Rebound – How quickly a system can recover from a failure
 - Graceful extensibility – Ability to react to things that aren't expected
 - Sustained adaptability – Ability to modify a system in response to new types of failures that are discovered

Scale and Resilience

Introduction

- How much failure is acceptable?
- Can your system still operate in a degraded state?
- How will failures be detected and responded to?

Scale and Resilience

Timeouts and Retries

- Base your timeouts and retries on your expected response times and latency
- Consider implementing an exponential backoff and possibly jitter

Scale and Resilience

Circuit Breakers

- Ensures a system can handle failures and maintain functionality in the presence of disruptions
 - Helps to prevent a cascading failure
- When a failure condition is detected, a circuit breaker can prevent additional calls to the failed component
- When a circuit breaker is "open", calling components can choose a different path or use cached data

Scale and Resilience

Bulkheads

- In a microservice architecture, a bulkhead another technique for preventing an issue from cascading through a system
- One common implementation is to place a request queue in front of a commonly used service
 - Can absorb an abnormal load
 - Can be monitored to trigger the deployment of additional resources

Scale and Resilience

Idempotency

- Means that performing the same action multiple times will have the same result as performing it just once
- Supports many of the other resiliency mechanisms
- Can be implemented with unique message IDs
 - Adding a time component can reduce overhead

Scale and Resilience

Vertical vs. Horizontal Scaling

- Vertical scaling involves adding more resources to an instance (vCPUs, RAM, etc.)
- Horizontal scaling involves adding more instances
 - More flexible and resilient but adds complexity

Scale and Resilience

Caching

- Can provide improved performance and resilience
- Useful when the resources used to retrieve the data from the cache are less than going directly to the source
- Useful for read heavy workloads and to support patterns such as circuit breaker
- System needs to be able to tolerate stale data

"There are only two hard things in Computer Science: cache invalidation and naming things" – Phil Karlton

Scale and Resilience

Chaos Engineering

- Proactively address potential failures by intentionally injecting controlled chaos
- Useful for testing both technical failure mechanisms as well as human response
- Netflix "chaos monkey" is an extreme example

Microservice Architecture with ASP.NET Core

Containerization

- dotnet publish
- Docker
- Orchestration
- Kubernetes

Containerization


dotnet publish

- The dotnet publish command compiles an app and copies the files required for deployment into a publish folder
 - Used by Visual Studio's [Build > Publish] wizard
- When performing a publish, the app can be...
 - Framework-dependent
 - Does not include the .NET runtime – correct version must already be present on the deployment machine
 - Self-contained
 - Does include the .NET runtime
 - Must choose the target architecture at publish time

Containerization

Kestrel


- The Kestrel web server is supported on all platforms where .NET is supported
- You can use Kestrel by itself or with a reverse proxy server



Containerization

IIS

- Out-of-process hosting



- In-process hosting



Containerization

Documentation

- There is extensive documentation for a wide variety of deployment scenarios
 - Apache httpd, Nginx, etc.

`docs.microsoft.com/en-us/aspnet/core/
host-and-deploy`

Containerization

Docker

- Docker is an open platform that enables developers and administrators to build images, ship, and run applications in a loosely isolated environment called a container

`www.docker.com`

Containerization

Docker

- Developer – "Helps me to eliminate the 'works on my machine' problem"
 - Application is executed and debugged in the same container that is deployed to production
- Administrator – "Allows me to treat hardware instances less like 'pets' and more like 'cattle'"
 - Hardware resources come-and-go with minimal configuration requirements – just need the ability to run a container

Containerization

Docker

- The Docker platform uses the Docker engine to build and package apps as Docker images
- Docker images are created using files written in the Dockerfile format

Containerization

Docker

- Microsoft provides a collection of official images to act as the starting point for your own images
 - Have the .NET Core runtime pre-installed
 - Some have the .NET Core SDK installed and can be used as a build server
- Available on Docker Hub
 - hub.docker.com/_/microsoft-dotnet-sdk
 - hub.docker.com/_/microsoft-dotnet-aspnet

Containerization

Docker

- A container is a running instance of an image
- When running an ASP.NET Core application in a container, it is necessary to map the internal container port to a port on the host machine

```
docker run -d -p 80:80 ecomm/api
```

- The Microsoft images are preconfigured (via an environment variable) to run Kestrel on port 80
- Documentation has extensive information about configuring Kestrel (HTTPS, HTTP/2, host filtering, etc.)

Containerization

Docker

- Visual Studio supports building, running, and debugging containerized ASP.NET Core applications
 - Must have Docker for Windows installed
- You can enable Docker support when creating a project
- You can also add Docker support to an existing app by selecting [Add > Docker Support]

Containerization

Docker

- Docker containers can run on any machine with the Docker engine installed
- Both Amazon AWS and Microsoft Azure have extensive support for hosting containers
 - AWS EC2 Container Service and Container Registry
 - Azure Container Service and Container Registry

Containerization

Orchestration

- "Automatically provisions, deploys, scales, and manages containerized applications without worrying about the underlying infrastructure"
- Scheduling, Affinity, Health monitoring, Failover, Scaling, Networking, Service discovery, Coordinated application upgrades

Containerization

Orchestration

- Every major cloud provider offers options for container orchestration
 - AWS ECS / EKS
 - Azure AKS
- Some are custom while others are built around Kubernetes

Containerization

Kubernetes

- Open-source container orchestration tool
- Originally developed by Google
- Components include...
 - Pod – Abstraction of containers (wrapper)
 - Service – Provides communication and load balancing
 - Ingress – Routes traffic into a cluster
 - ConfigMap / Secret – External configuration
 - Volume – Data persistence
 - Deployment – Replication of pods

Containerization

Kubernetes

- Minikube can be used to run an entire cluster locally

Microservice Architecture with ASP.NET Core

Security

- Authentication
- Authorization
- Web API Authentication
- Secrets Management

Security

Authentication

- Authentication is the process of discovering the identity of an entity through an identifier and verifying the identity by validating the credentials provided by the entity
- It is common to validate the credentials once and generate a temporary unique token that is used to authorize other actions
 - The token can contain other information (claims) about the entity including group membership and privileges
 - Care must be taken to ensure the token is sent in a secure manner and the integrity of the token can be verified

Security

Authentication

- The authentication middleware is added by calling `UseAuthentication` in the `Configure` method
- If a request is not authenticated, authorization invokes a challenge using the default (or specified) authentication scheme
- Authentication challenge examples:
 - Redirection to a login page for a Web UI
 - 401 response with `www-authentication: bearer` header for a Web API

Security

Authorization

- The `Authorize` attribute can be used to authorize access to specific functionality
- Can be applied at the action or controller level
- If nothing else is specified, the attribute simply ensures the current user has been authenticated
- The `AllowAnonymous` attribute can be used to opt-out an action

```
[Authorize]
public class AdminController : Controller
{
    public IActionResult Dashboard() { ... }

    [AllowAnonymous]
    public IActionResult Login() { ... }
}
```

Security

Authorization

- With claims-based authorization, you can define a policy that specifies what claims must be present for a user to be authorized

```
services.AddAuthorization(options => {  
    options.AddPolicy("AdminsOnly", policy =>  
        policy.RequireClaim(ClaimTypes.Role, "Admin"));  
});
```

- The Authorize attribute can then be used to enforce the policy

```
[Authorize(Policy="AdminsOnly")]  
public IActionResult Dashboard() { ... }
```

Security

Authorization

- Programmatic authorization checks can also be performed within an action to enforce function-level access control

```
public IActionResult Delete(int id)  
{  
    if (User.HasClaim("SomeClaim")) {  
        //  
    }  
}
```

```
public IActionResult Dashboard()  
{  
    foreach (var claim in User.Claims)  
    {  
        // decide what the user should see  
    }  
}
```

Security

Web API Authentication

- For API authentication, cookie-based authentication is typically not used, and it is not possible to directly present a UI to gather credentials
- In some scenarios, a simple form of authentication where the credentials are passed with each request is sufficient
 - HTTP Basic
 - Pre-shared key (PSK) authentication
- In a microservice architecture, an API Gateway can be used to perform authentication and proxy requests to other services

Security

Web API Authentication

- A modern flexible approach to API authentication is to use bearer token authentication
 - ASP.NET Core supports OAuth 2.0 and OpenID Connect

```
services.AddAuthentication(options => {
    options.DefaultAuthenticateScheme =.AspNetCore.Authentication
        .JwtBearer.JwtBearerDefaults.AuthenticationScheme;

    options.DefaultChallengeScheme =.AspNetCore.Authentication
        .JwtBearer.JwtBearerDefaults.AuthenticationScheme;
})
.AddJwtBearer(options => {
    options.Authority = identityUrl;
    options.Audience = "EComm";
});
```

Security

Web API Authentication

- Authority property is the address of the token-issuing authentication server
 - Used by the middleware to obtain the public key for validation of the token signature
- Audience property represents the resource the token grants access to
 - Value must match the parameter in the token

Security

Web API Authentication

- The token issuing server can be an external identity provider (Facebook, Twitter, Microsoft, etc.) or you can create your own using ASP.NET Core
- If issuing your own security tokens, it is recommended to use a third-party library to handle many of the security-related details
 - IdentityServer4
 - OpenIddict

Security

Web API Authentication

- It is also possible to create a custom authentication scheme / handler
 - Can create a subclass `AuthenticationHandler` and override `HandleAuthenticateAsync`
- When developing a custom authentication scheme, extreme care should be taken to ensure security vulnerabilities are not introduced

Security

Secrets Management

- Secrets are configuration values that are in some way sensitive
 - Connection strings
 - API keys
- It is a good practice for secrets not to be in source code or config files (things that end up in source control)
- Secrets should be made available to the production environment through a controlled means
 - Environment variables
 - Azure Key Vault
 - AWS Secrets Manager

Security

Secrets Management

- For local development, the Secret Manager tool can be used
- Secrets are stored in a JSON configuration file in a system-protected user profile folder on the local machine
- Use the init command or select Manage User Secrets in Visual Studio to enable user secrets for a project

```
dotnet user-secrets init
```

- Adds an entry into the project file

```
<PropertyGroup>  
  <TargetFramework>netcoreapp3.1</TargetFramework>  
  <UserSecretsId>79a3edd0-2092-40a2-a04d-dcb46d5ca9ed</UserSecretsId>  
</PropertyGroup>
```

Security

Secrets Management

- Secrets can be added via the command-line or by using the Manage User Secrets command in Visual Studio

```
dotnet user-secrets set "Movies:ServiceApiKey" "12345"
```

- CreateDefaultBuilder calls AddUserSecrets when the EnvironmentName is Development